

NAAN MUDHALVAN

FINAL DEVELOPMENT:

diabetes_predictions:

The **diabetes_prediction_dataset.csv** file contains medical and demographic data of patients along with their diabetes status, whether positive or negative. It consists of various features such as age, gender, body mass index (BMI), hypertension, heart disease, smoking history, HbA1c level, and blood glucose level. The Dataset can be utilized to construct machine learning models that can predict the likelihood of diabetes in patients based on their medical history and demographic details.

Prediction of diabetes

In [1]:

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn import linear_model
from sklearn.metrics import r2_score, confusion_matrix

# import warnings
import warnings
# filter warnings
warnings.filterwarnings('ignore')

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will
# list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
```

```
print(os.path.join(dirname, filename))
```

You can write up to 20GB to the current directory (/kaggle/working/) that gets preserved as output when you create a version using "Save & Run All"

You can also write temporary files to /kaggle/temp/, but they won't be saved outside of the current session

```
/kaggle/input/diabetes-prediction-dataset/diabetes_prediction_dataset.csv
```

Load and Check Data

In [2]:

```
data =  
pd.read_csv("/kaggle/input/diabetes-prediction-dataset/diabetes_prediction_dataset.csv")
```

In [3]:

```
data.head()
```

Out[3]:

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level	diabetes
0	Female	80.0	0	1	never	25.19	6.6	140	0

1	Female	54.0	0	0	No Info	27.32	6.6	80	0
2	Male	28.0	0	0	never	27.32	5.7	158	0
3	Female	36.0	0	0	current	23.45	5.0	155	0
4	Male	76.0	1	1	current	20.14	4.8	155	0

In [4]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                 100000 non-null object
1   age                   100000 non-null float64
2   hypertension           100000 non-null int64
3   heart_disease          100000 non-null int64
4   smoking_history        100000 non-null object
5   bmi                   100000 non-null float64
6   HbA1c_level            100000 non-null float64
7   blood_glucose_level    100000 non-null int64
8   diabetes               100000 non-null int64
dtypes: float64(3), int64(4), object(2)
memory usage: 6.9+ MB
```

In [5]:

```
def bar_plot(variable):
    '''
```

```

input: variable ex: "Sex"
output: bar plot & value count
'''
# get feature
var = data[variable]
# count number of categorical variable(value/sample)
var_value = var.value_counts()
# visualize
plt.figure(figsize=(6,3))
plt.bar(var_value.index, var_value,width= 1/(var.unique().size))
plt.xticks(var_value.index, var_value.index.values)
plt.ylabel("Frequency")
plt.title(variable)
plt.show()
print("{}: \n {}".format(variable,var_value))


```

In [6]:

```

category1 = ["hypertension", "heart_disease", "smoking_history"]
for c in category1:
    bar_plot(c)


```



```

hypertension:
hypertension
0      92515
1       7485
Name: count, dtype: int64

```




```

heart_disease:
heart_disease
0      96058

```

```
1      3942
Name: count, dtype: int64
```



```
smoking_history:
  smoking_history
No Info      35816
never        35095
former        9352
current       9286
not current   6447
ever          4004
Name: count, dtype: int64
```

Change the Data Type

Making data types integer

```
In [7]:
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   gender              100000 non-null object  
 1   age                 100000 non-null float64
 2   hypertension        100000 non-null int64   
 3   heart_disease       100000 non-null int64   
 4   smoking_history     100000 non-null object
```

```
5   bmi                100000 non-null   float64
6   HbA1c_level        100000 non-null   float64
7   blood_glucose_level 100000 non-null   int64
8   diabetes           100000 non-null   int64
dtypes: float64(3), int64(4), object(2)
memory usage: 6.9+ MB
```

```
In [8]:
pd.unique(data.smoking_history)
```

```
Out[8]:
array(['never', 'No Info', 'current', 'former', 'ever', 'not current'],
      dtype=object)
```

```
In [9]:
pd.unique(data.gender)
```

```
Out[9]:
array(['Female', 'Male', 'Other'], dtype=object)
```

For changing objects to integer

```
In [10]:
def change_string_to_int(column):
    variables=pd.unique(data[column])
    for item in range(variables.size):
        data[column]=[item if each==variables[item] else each for each in
data[column]]
    return data[column]
```

```
In [11]:
data["gender"]=change_string_to_int("gender")
```

```
In [12]:
```

```
data["smoking_history"]=change_string_to_int("smoking_history")
```

```
In [13]:
```

```
data.head()
```

```
Out[13]:
```

	gender	age	hyperte nsion	heart_di sease	smoking _history	bmi	HbA1c_ level	blood_g lucose_ level	diabete s
0	0	80.0	0	1	0	25.19	6.6	140	0
1	0	54.0	0	0	1	27.32	6.6	80	0
2	1	28.0	0	0	0	27.32	5.7	158	0
3	0	36.0	0	0	2	23.45	5.0	155	0
4	1	76.0	1	1	2	20.14	4.8	155	0

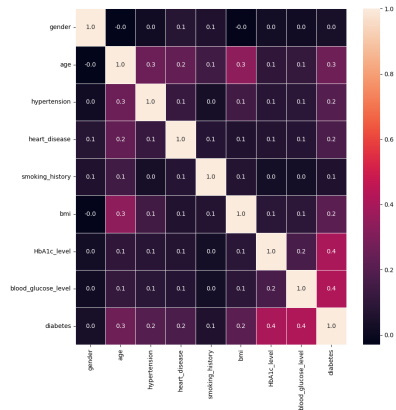
```
In [14]:
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   gender                100000 non-null  int64   
 1   age                   100000 non-null  float64  
 2   hypertension          100000 non-null  int64   
 3   heart_disease         100000 non-null  int64   
 4   smoking_history       100000 non-null  int64   
 5   bmi                   100000 non-null  float64  
 6   HbA1c_level           100000 non-null  float64  
 7   blood_glucose_level   100000 non-null  int64   
 8   diabetes              100000 non-null  int64   
dtypes: float64(3), int64(6)
memory usage: 6.9 MB
```

Removing from the Data

```
In [15]:
f,ax = plt.subplots(figsize=(10, 10))
sns.heatmap(data.corr(numeric_only=True), annot=True, linewidths=.5, fmt=
'.1f',ax=ax,)
plt.show()
```

```
In [16]:
data.drop("gender",axis=1,inplace=True)
```

```
In [17]:
data.head()
```


Out[17]:

	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level	diabetes
0	80.0	0	1	0	25.19	6.6	140	0
1	54.0	0	0	1	27.32	6.6	80	0
2	28.0	0	0	0	27.32	5.7	158	0

3	36.0	0	0	2	23.45	5.0	155	0
4	76.0	1	1	2	20.14	4.8	155	0

Transform the Data

```
In [18]:  
data.describe()
```



Out[18]:

	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level	diabetes
count	100000.00000	100000.0000	100000.00000	100000.00000	100000.00000	100000.00000	100000.00000	100000.00000
mean	41.885856	0.07485	0.039420	1.306950	27.320767	5.527507	138.058060	0.085000
std	22.516840	0.26315	0.194593	1.454501	6.636783	1.070672	40.708136	0.278883

min	0.080000	0.00000	0.000000	0.000000	10.01000 0	3.500000	80.00000 0	0.000000
25%	24.00000 0	0.00000	0.000000	0.000000	23.63000 0	4.800000	100.0000 00	0.000000
50%	43.00000 0	0.00000	0.000000	1.000000	27.32000 0	5.800000	140.0000 00	0.000000
75%	60.00000 0	0.00000	0.000000	2.000000	29.58000 0	6.200000	159.0000 00	0.000000
max	80.00000 0	1.00000	1.000000	5.000000	95.69000 0	9.000000	300.0000 00	1.000000

Normalization

In [19]:

```
data = (data - data.min()) / (data.max() - data.min())
```

In [20]:

```
data.head()
```

Out[20]:

	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level	diabetes
0	1.000000	0.0	1.0	0.0	0.177171	0.563636	0.272727	0.0
1	0.674675	0.0	0.0	0.2	0.202031	0.563636	0.000000	0.0
2	0.349349	0.0	0.0	0.0	0.202031	0.400000	0.354545	0.0
3	0.449449	0.0	0.0	0.4	0.156863	0.272727	0.340909	0.0
4	0.949950	1.0	1.0	0.4	0.118231	0.236364	0.340909	0.0

In [21]:

```
data["hypertension"] = data["hypertension"].astype("int64")
data["heart_disease"] = data["heart_disease"].astype("int64")
data["diabetes"] = data["diabetes"].astype("int64")
```

In [22]:

```
data.head()
```

Out[22]:

	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level	diabetes
0	1.000000	0	1	0.0	0.177171	0.563636	0.272727	0
1	0.674675	0	0	0.2	0.202031	0.563636	0.000000	0
2	0.349349	0	0	0.0	0.202031	0.400000	0.354545	0
3	0.449449	0	0	0.4	0.156863	0.272727	0.340909	0
4	0.949950	1	1	0.4	0.118231	0.236364	0.340909	0

Split and Train

In [23]:

```
x = data.drop("diabetes",axis=1)
y = data.diabetes
```

In [24]:

```
x_train, x_test, y_train, y_test =
train_test_split(x,y,test_size=0.3,random_state=42)
```

```
In [25]:  
y_train = np.array(y_train).reshape(-1,1)  
y_test = np.array(y_test).reshape(-1,1)
```

```
In [26]:  
print("x train: ",x_train.shape)  
print("x test: ",x_test.shape)  
print("y train: ",y_train.shape)  
print("y test: ",y_test.shape)
```

```
x train: (70000, 7)  
x test: (30000, 7)  
y train: (70000, 1)  
y test: (30000, 1)
```

```
In [27]:  
logreg = linear_model.LogisticRegression(random_state = 42,max_iter= 200,)  
print("test accuracy: {}".format(logreg.fit(x_train,  
y_train).score(x_test, y_test)))  
print("train accuracy: {}".format(logreg.fit(x_train,  
y_train).score(x_train, y_train)))
```

```
test accuracy: 0.9587666666666667  
train accuracy: 0.9610571428571428
```

```
In [28]:  
y_pred = logreg.fit(x_train, y_train).predict(x_test)  
cm = confusion_matrix(y_test,y_pred)
```

```
In [29]:
```

```
ax= plt.subplot()  
sns.heatmap(cm, annot=True, fmt='g', ax=ax);  
ax.set_xlabel('Predicted labels');ax.set_ylabel('True labels');  
ax.set_title('Confusion Matrix');  
ax.xaxis.set_ticklabels(['diabetes', 'not_diabetes']);  
ax.yaxis.set_ticklabels(['diabetes', 'not_diabetes']);
```

